

CMDA-4654

Project 1

Spencer Paragas and Christopher Mascis

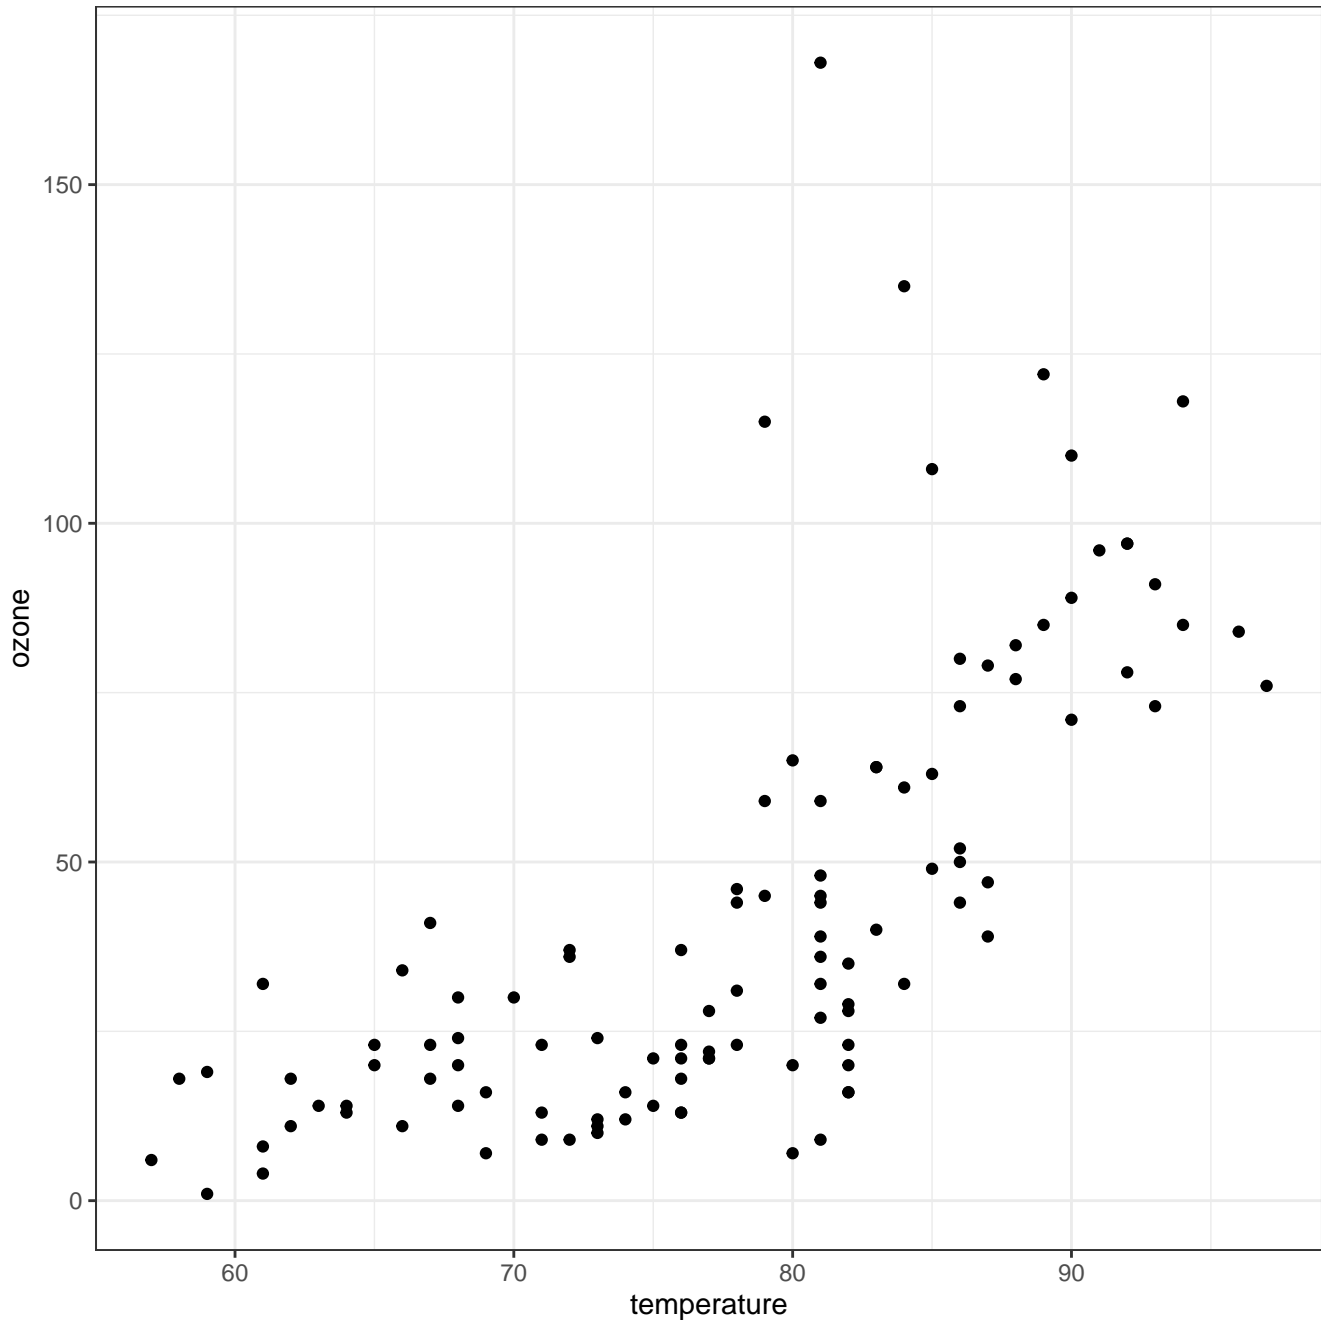
10/19/2021

Problem 1

```
library(dplyr)
library(ggplot2)

setwd("C:/Users/Spencer/Downloads")
load("ozone.RData")

ggplot(ozone, aes(x = temperature, y = ozone)) + theme_bw() + geom_point()
```



Part a

Fitting polynomials of different degrees

```
model1 <- lm(ozone ~ poly(radiation,1), data = ozone)
model2 <- lm(ozone ~ poly(radiation,2), data = ozone)
model3 <- lm(ozone ~ poly(radiation,3), data = ozone)
model4 <- lm(ozone ~ poly(radiation,4), data = ozone)
```

```

model5 <- lm(ozone ~ poly(radiation,5), data = ozone)
model6 <- lm(ozone ~ poly(radiation,6), data = ozone)

summary(model1)

Call:
lm(formula = ozone ~ poly(radiation, 1), data = ozone)

Residuals:
    Min       1Q   Median       3Q      Max
-48.292 -21.361  -8.864  16.373 119.136

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      42.099      2.974   14.15 < 2e-16 ***
poly(radiation, 1) 121.572     31.335    3.88 0.000179 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 31.33 on 109 degrees of freedom
Multiple R-squared:  0.1213,    Adjusted R-squared:  0.1133
F-statistic: 15.05 on 1 and 109 DF,  p-value: 0.0001793

summary(model2)

Call:
lm(formula = ozone ~ poly(radiation, 2), data = ozone)

Residuals:
    Min       1Q   Median       3Q      Max
-40.155 -22.793  -6.438  18.061 115.117

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      42.099      2.832   14.864 < 2e-16 ***
poly(radiation, 2)1 121.572     29.840    4.074 8.84e-05 ***
poly(radiation, 2)2 -104.178     29.840   -3.491 0.000698 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 29.84 on 108 degrees of freedom
Multiple R-squared:  0.2104,    Adjusted R-squared:  0.1958
F-statistic: 14.39 on 2 and 108 DF,  p-value: 2.875e-06

summary(model3)

Call:
lm(formula = ozone ~ poly(radiation, 3), data = ozone)

Residuals:
    Min       1Q   Median       3Q      Max
-45.05 -19.44  -4.52  15.61 109.95

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      42.10      2.78   15.141 < 2e-16 ***
poly(radiation, 3)1 121.57     29.29    4.150 6.69e-05 ***
poly(radiation, 3)2 -104.18     29.29   -3.556 0.000562 ***
poly(radiation, 3)3  -65.93     29.29   -2.251 0.026450 *

```

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 29.29 on 107 degrees of freedom
Multiple R-squared:  0.2461,    Adjusted R-squared:  0.225
F-statistic: 11.65 on 3 and 107 DF,  p-value: 1.154e-06
```

```
summary(model4)
```

```
Call:
lm(formula = ozone ~ poly(radiation, 4), data = ozone)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-44.929 -19.519  -4.647   15.700  110.071
```

```
Coefficients:
                Estimate Std. Error t value Pr(>|t|)
(Intercept)         42.099      2.793   15.072 < 2e-16 ***
poly(radiation, 4)1  121.572     29.428    4.131 7.23e-05 ***
poly(radiation, 4)2 -104.178     29.428   -3.540 0.000596 ***
poly(radiation, 4)3  -65.932     29.428   -2.240 0.027149 *
poly(radiation, 4)4    4.777     29.428    0.162 0.871356
```

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 29.43 on 106 degrees of freedom
Multiple R-squared:  0.2463,    Adjusted R-squared:  0.2179
F-statistic: 8.661 on 4 and 106 DF,  p-value: 4.352e-06
```

```
summary(model5)
```

```
Call:
lm(formula = ozone ~ poly(radiation, 5), data = ozone)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-46.076 -16.978  -5.621   16.472  108.924
```

```
Coefficients:
                Estimate Std. Error t value Pr(>|t|)
(Intercept)         42.099      2.801   15.029 < 2e-16 ***
poly(radiation, 5)1  121.572     29.511    4.119 7.59e-05 ***
poly(radiation, 5)2 -104.178     29.511   -3.530 0.000618 ***
poly(radiation, 5)3  -65.932     29.511   -2.234 0.027595 *
poly(radiation, 5)4    4.777     29.511    0.162 0.871718
poly(radiation, 5)5   18.759     29.511    0.636 0.526389
```

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 29.51 on 105 degrees of freedom
Multiple R-squared:  0.2492,    Adjusted R-squared:  0.2135
F-statistic: 6.971 on 5 and 105 DF,  p-value: 1.168e-05
```

```
summary(model6)
```

```
Call:
lm(formula = ozone ~ poly(radiation, 6), data = ozone)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-45.079	-17.204	-5.404	16.483	109.921

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	42.099	2.812	14.971	< 2e-16 ***
poly(radiation, 6)1	121.572	29.626	4.104	8.11e-05 ***
poly(radiation, 6)2	-104.178	29.626	-3.516	0.00065 ***
poly(radiation, 6)3	-65.932	29.626	-2.225	0.02821 *
poly(radiation, 6)4	4.777	29.626	0.161	0.87221
poly(radiation, 6)5	18.759	29.626	0.633	0.52800
poly(radiation, 6)6	-12.854	29.626	-0.434	0.66527

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 29.63 on 104 degrees of freedom

Multiple R-squared: 0.2506, Adjusted R-squared: 0.2073

F-statistic: 5.795 on 6 and 104 DF, p-value: 3.079e-05

The polynomial fit that appears to work the best is the one with degree 3.

Part b

Writing the function that carries out LOESS regression

```
myloess <- function(x, y, span = 0.5, degree = 1, show.plot = TRUE) {

  # Getting range of values
  xrange <- diff(range(x))

  # Checking span meets requirements and setting width
  if (between(span, 0, 1)) {
    width <- span*xrange
  }
  else {
    stop("Span must be between 0 and 1 non-inclusive")
  }

  # Getting total number of points and windows (they'll be the same)
  N_total <- length(x)
  Win_total <- length(x)

  # Allocating space for vector of each window's population
  n_points <- vector(mode = "integer", length = length(x))

  # Allocating space for vector of fitted values
  yhat <- vector(mode = "numeric", length = length(x))

  # Combining our variables in data frame
  mydata <- cbind.data.frame(x, y)

  # Fitting each point
  for(x0 in x) {

    # Setting population of window
    sample <- subset(mydata, between(x, x0-width/2, x0+width/2))
    n <- length(sample[,1])

    # Getting weights into diagonal matrix
    weights <- (1-(abs(sample[,1] - x0)/width*2)^3)^3
    W_mat <- diag(weights, n, n)
```

```

# Checking degree and completing our regression accordingly
if (degree == 1) {

  X_mat <- cbind(rep(1, n), sample[,1])
  betahat <- solve( t(X_mat)%*%W_mat%*%X_mat ) %*% t(X_mat) %*% W_mat %*% sample[,2]

  # Getting fitted value
  yhat[which(x0 == x)] <- betahat[1] + betahat[2]*x0
}
else if (degree == 2) {

  X_mat <- cbind(rep(1, n), sample[,1], sample[,1]^2)
  betahat <- solve( t(X_mat)%*%W_mat%*%X_mat ) %*% t(X_mat) %*% W_mat %*% sample[,2]

  # Getting fitted value
  yhat[which(x0 == x)] <- betahat[1] + betahat[2]*x0 + betahat[3]*x0^2
}
else {
  stop("Degree must be set to either 1 or 2")
}

# Getting population of window
n_points[which(x0 == x)] <- n
}

# Calculating SSE, MSE, and residual standard error
SSE <- sum((y-yhat)^{2})
MSE <- SSE/(N_total-2)
RSE <- sqrt(MSE)

# Creating plot of final fit
loessplot <- ggplot(mydata, aes(x, y)) +
  geom_point(size = 3, alpha = 0.5, color = "grey") +
  geom_line(aes(x, yhat), color = "red", lty = 1) +
  xlab(deparse(substitute(x))) + ylab(deparse(substitute(y))) +
  ggtitle(paste("LOESS with degree =", degree,"and span =", span, sep = " "))

# Checking whether to show plot or not
if (show.plot == T) {
  print(loessplot)
}

# Returning named list
return(invisible(list(span = span, degree = degree, N_total = N_total, Win_total = Win_total,
  n_points = n_points, SSE = SSE, RSE = RSE, loessplot = loessplot)))
}

Determining LOESS regression fits on the data

# Creating an empty data frame
fit_table <- data.frame()

# Determining fits and putting info into data frame
for(j in 1:2) {
  for (i in seq(0.25, 0.75, by = 0.05)) {
    fit_table <- rbind(fit_table, c(i, j, myloess(ozone$temperature, ozone$ozone,
      span = i, degree = j, show.plot = F)$RSE))
  }
}

```

```
# Changing column names
colnames(fit_table) <- c("Span", "Degree", "RSE")
```

```
# Displaying data frame
fit_table
```

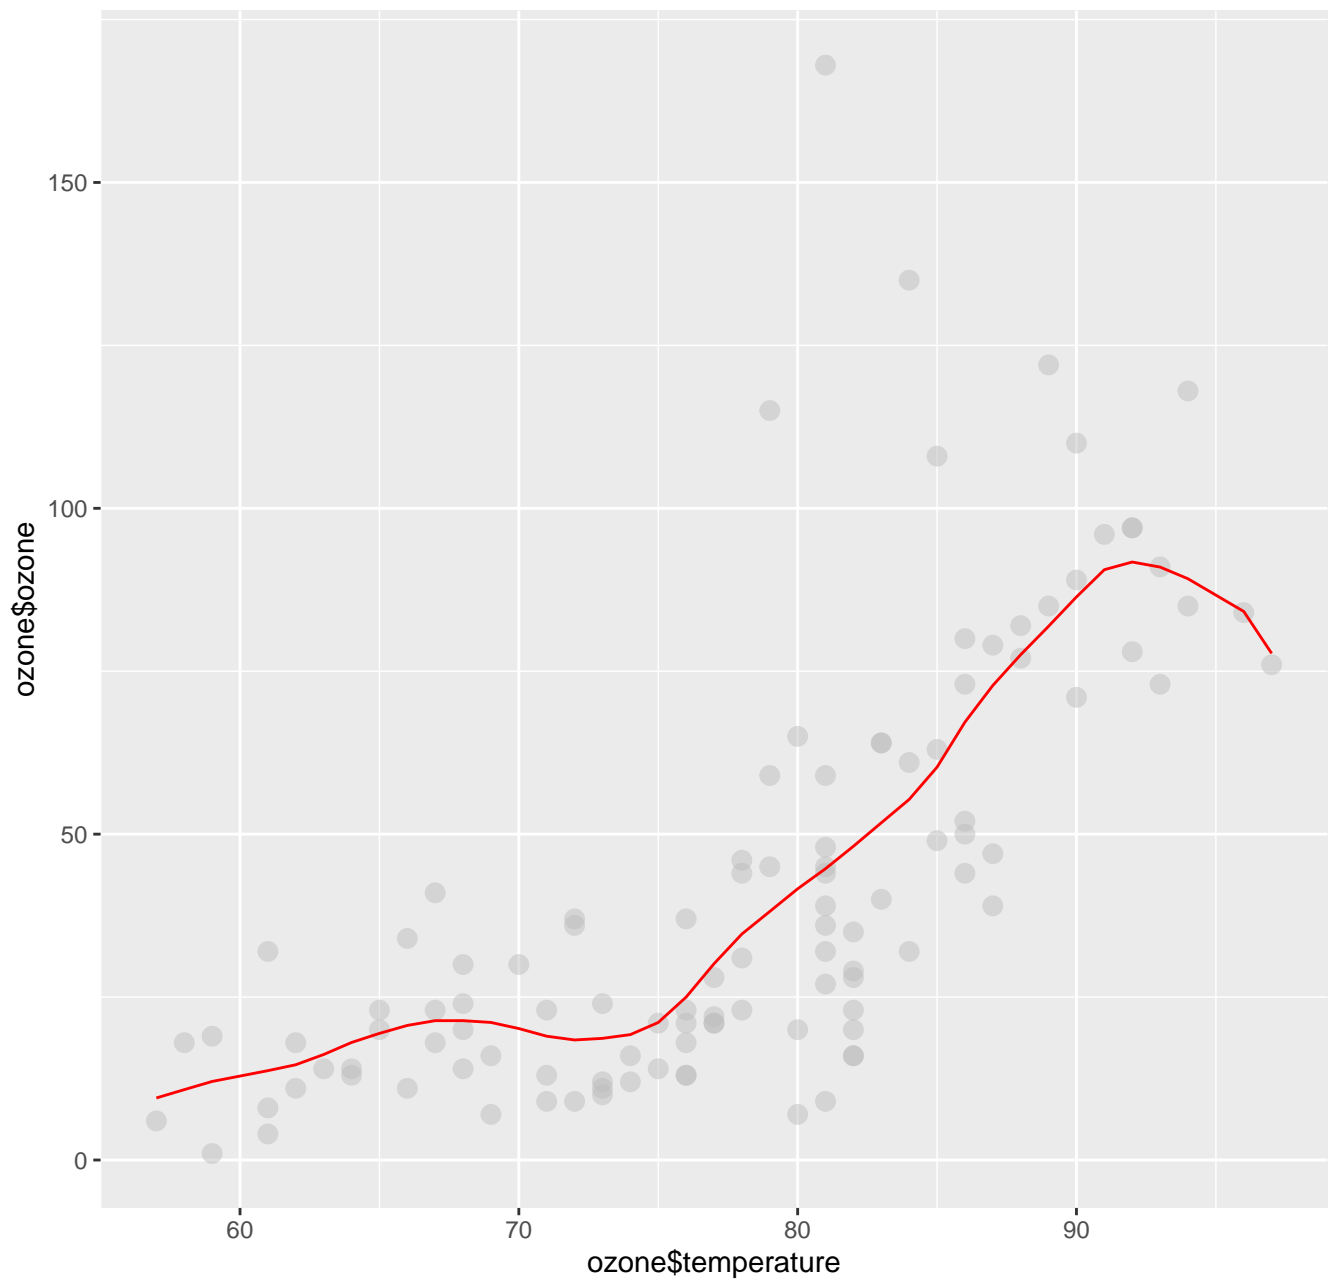
	Span	Degree	RSE
1	0.25	1	21.69723
2	0.30	1	21.84914
3	0.35	1	21.87405
4	0.40	1	21.87005
5	0.45	1	21.88484
6	0.50	1	21.92533
7	0.55	1	21.98316
8	0.60	1	22.05339
9	0.65	1	22.12531
10	0.70	1	22.18786
11	0.75	1	22.24218
12	0.25	2	20.70123
13	0.30	2	21.21788
14	0.35	2	21.65352
15	0.40	2	21.83486
16	0.45	2	21.82088
17	0.50	2	21.75230
18	0.55	2	21.71153
19	0.60	2	21.71786
20	0.65	2	21.76031
21	0.70	2	21.81288
22	0.75	2	21.86516

The three “best” fits with degree 1 appear to be span 0.25, 0.30, and 0.40. The three “best” fits with degree 2 appear to be span 0.25, 0.30, and 0.35.

Plotting the best fits found above

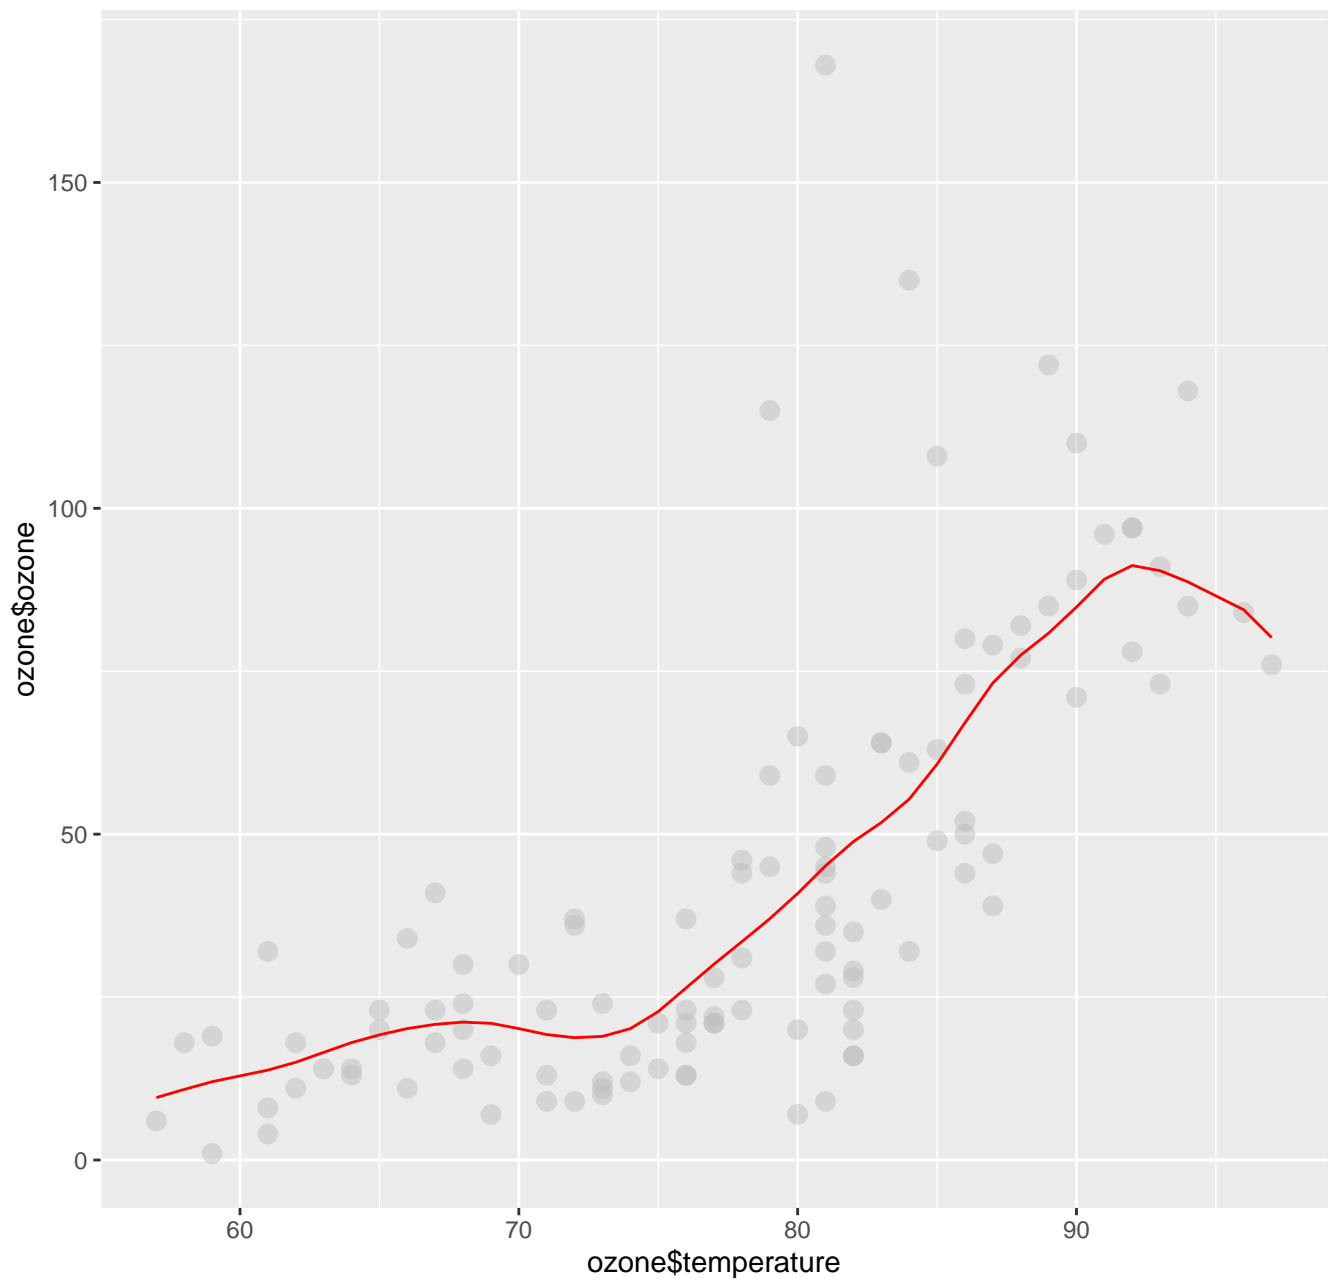
```
myloess(ozone$temperature, ozone$ozone,
span = 0.25, degree = 1, show.plot = F)$loessplot
```

LOESS with degree = 1 and span = 0.25



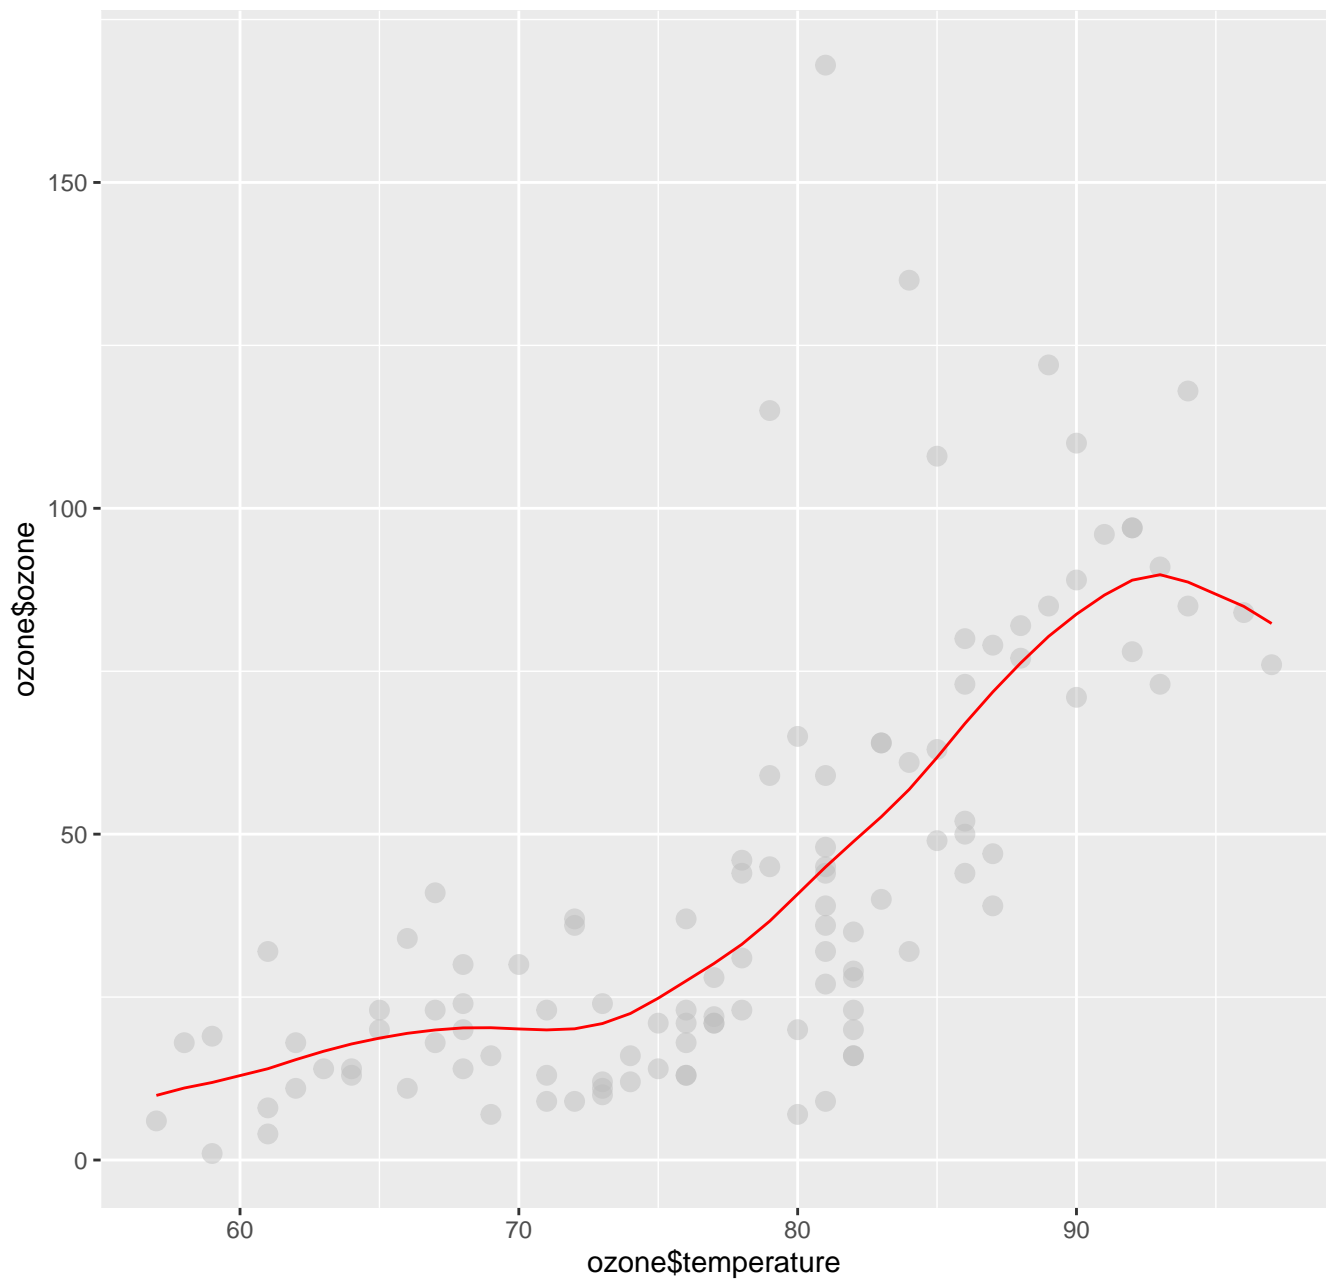
```
myloess(ozone$temperature, ozone$ozone,  
span = 0.30, degree = 1, show.plot = F)$loessplot
```


LOESS with degree = 1 and span = 0.3



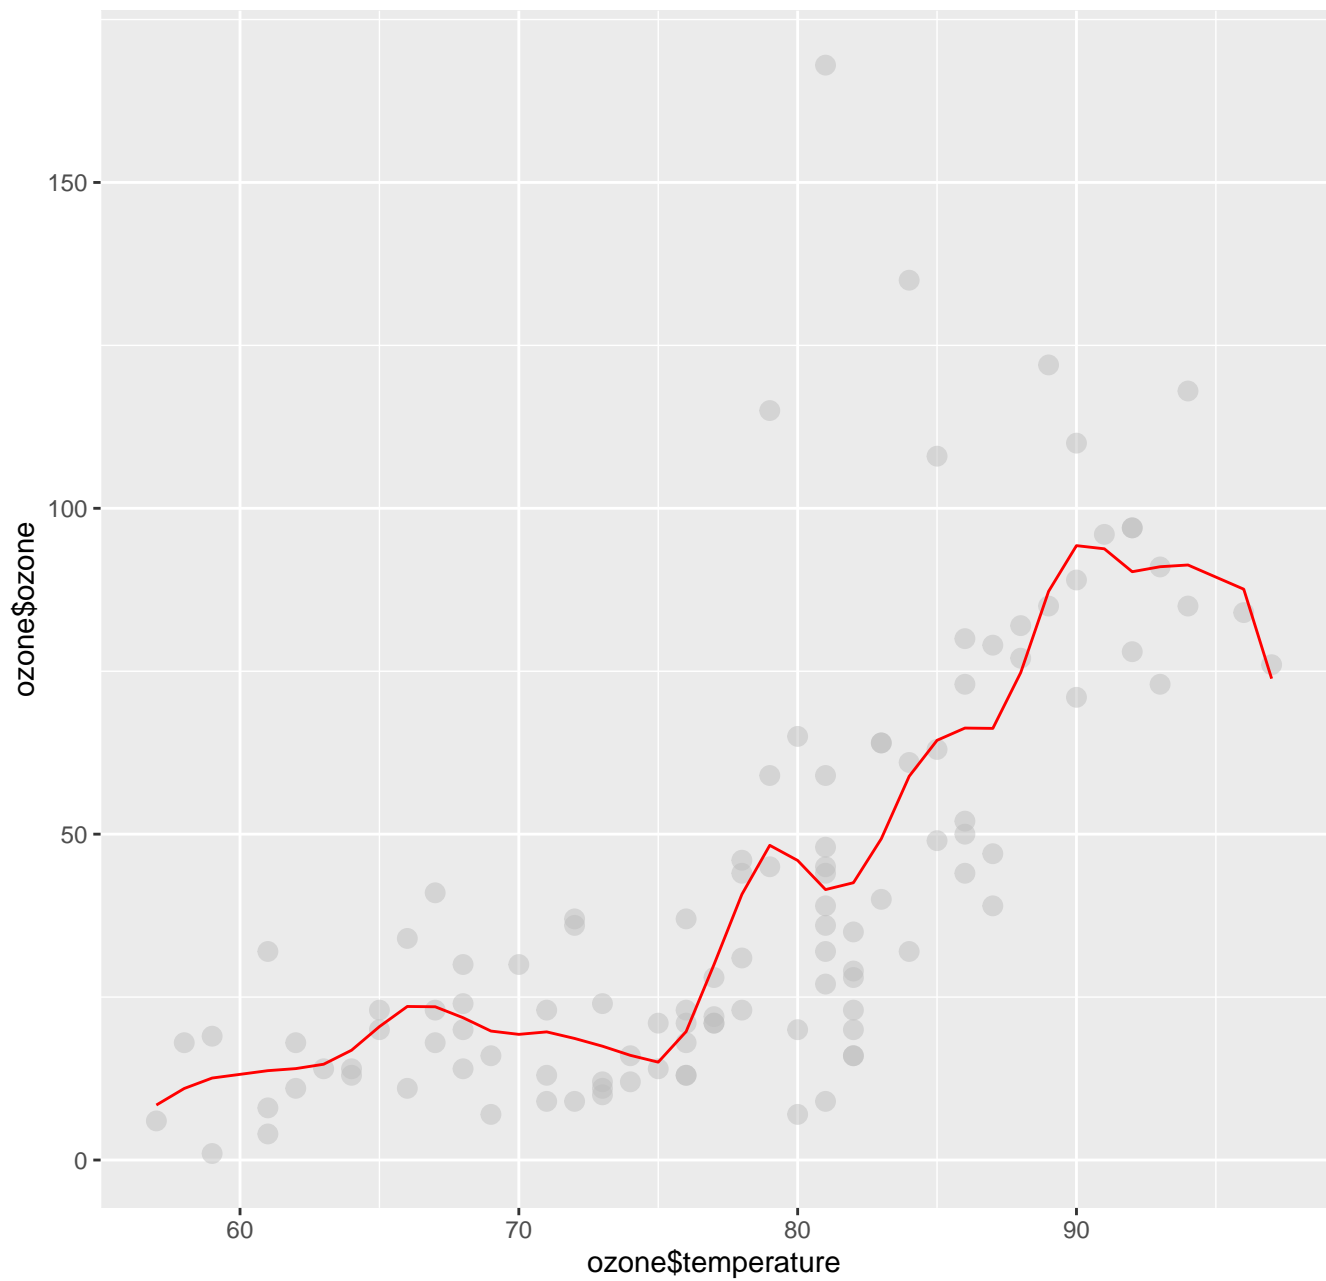
```
myloess(ozone$temperature, ozone$ozone,  
span = 0.40, degree = 1, show.plot = F)$loessplot
```

LOESS with degree = 1 and span = 0.4



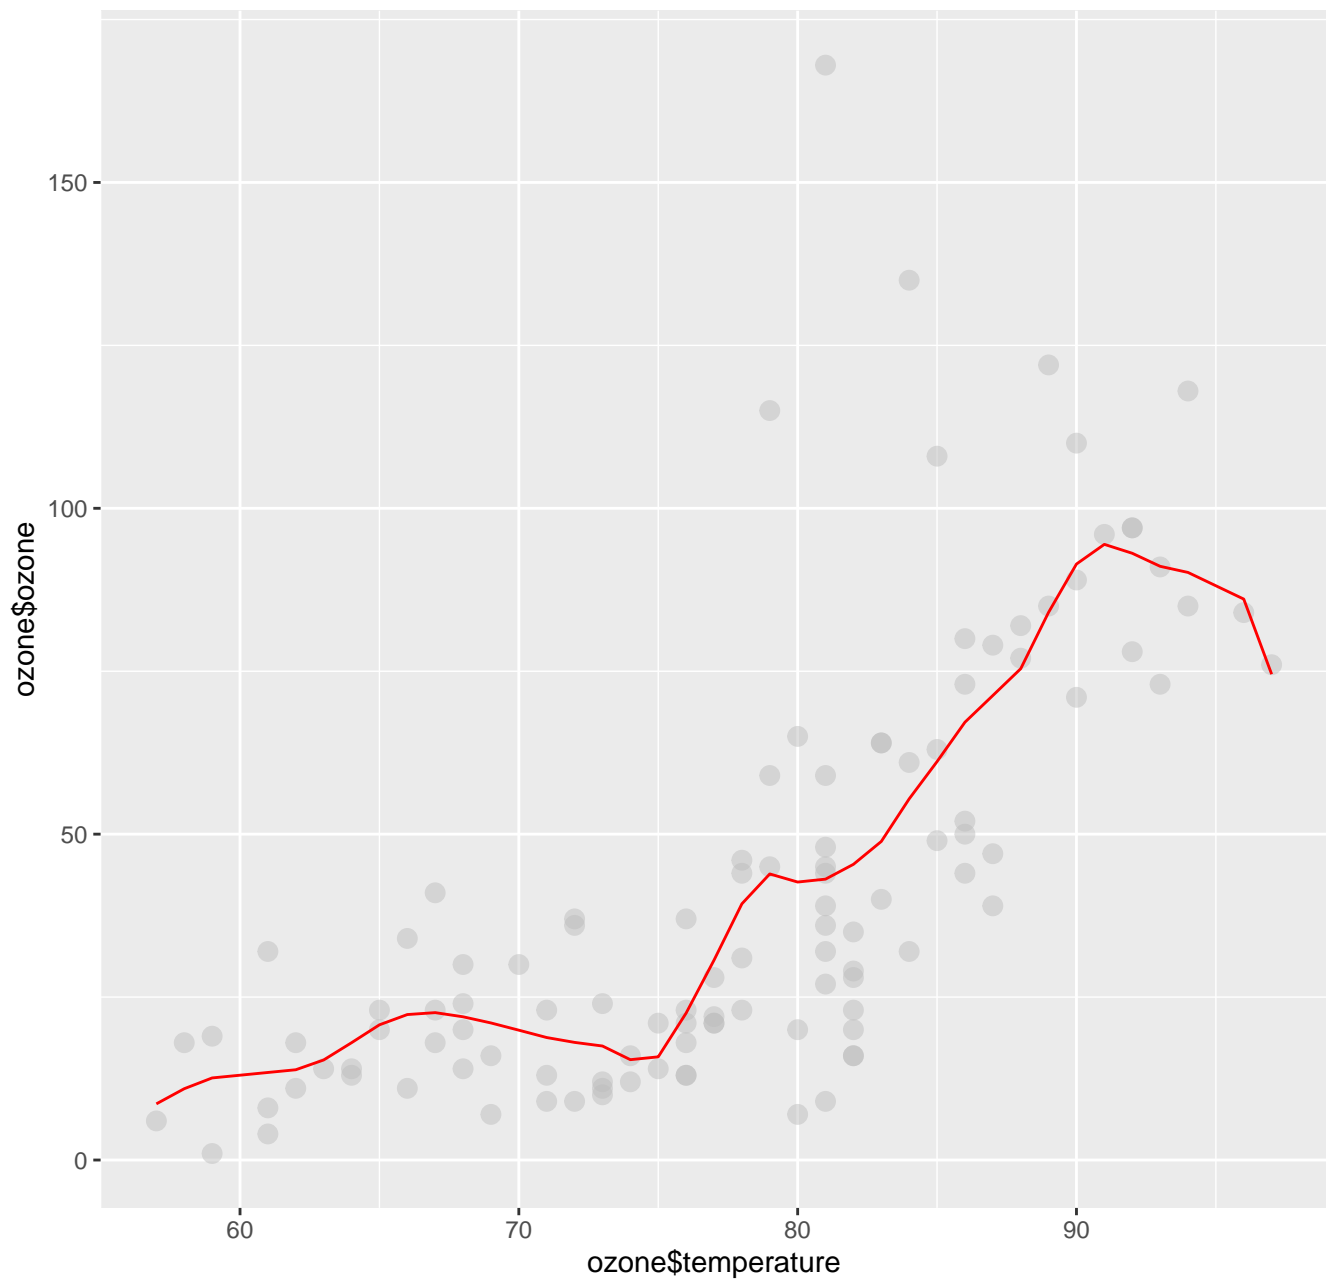
```
myloess(ozone$temperature, ozone$ozone,  
span = 0.25, degree = 2, show.plot = F)$loessplot
```

LOESS with degree = 2 and span = 0.25



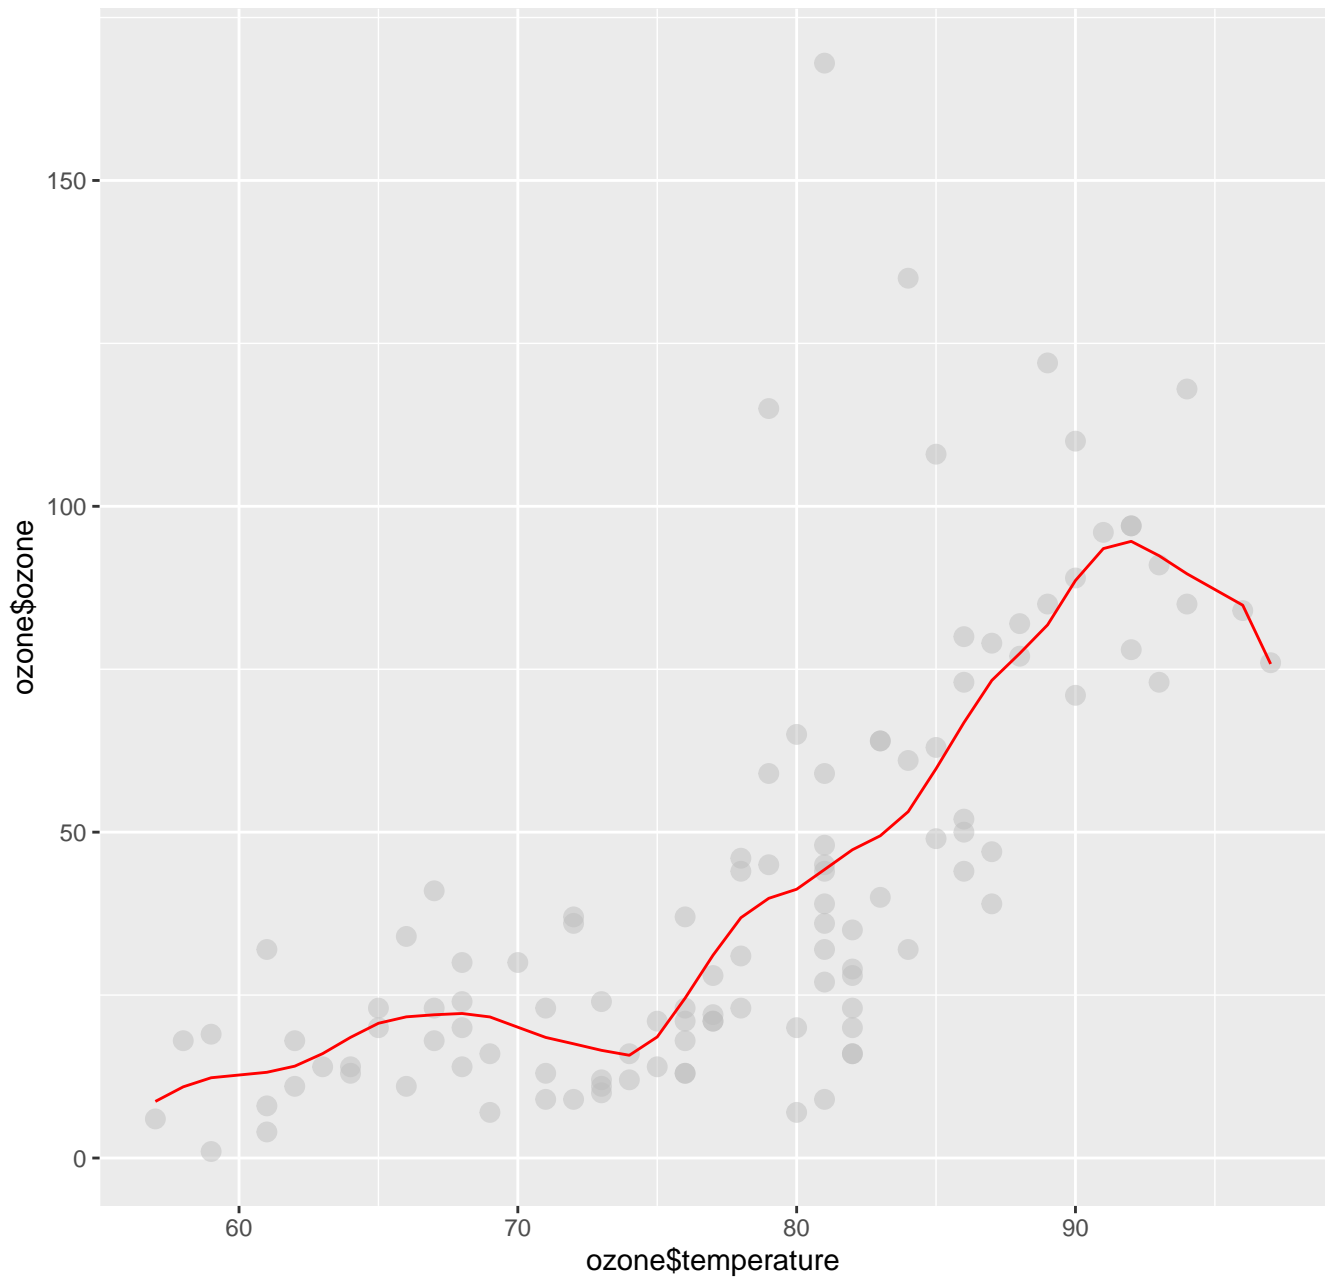
```
myloess(ozone$temperature, ozone$ozone,  
span = 0.30, degree = 2, show.plot = F)$loessplot
```

LOESS with degree = 2 and span = 0.3



```
myloess(ozone$temperature, ozone$ozone,  
span = 0.35, degree = 2, show.plot = F)$loessplot
```

LOESS with degree = 2 and span = 0.35

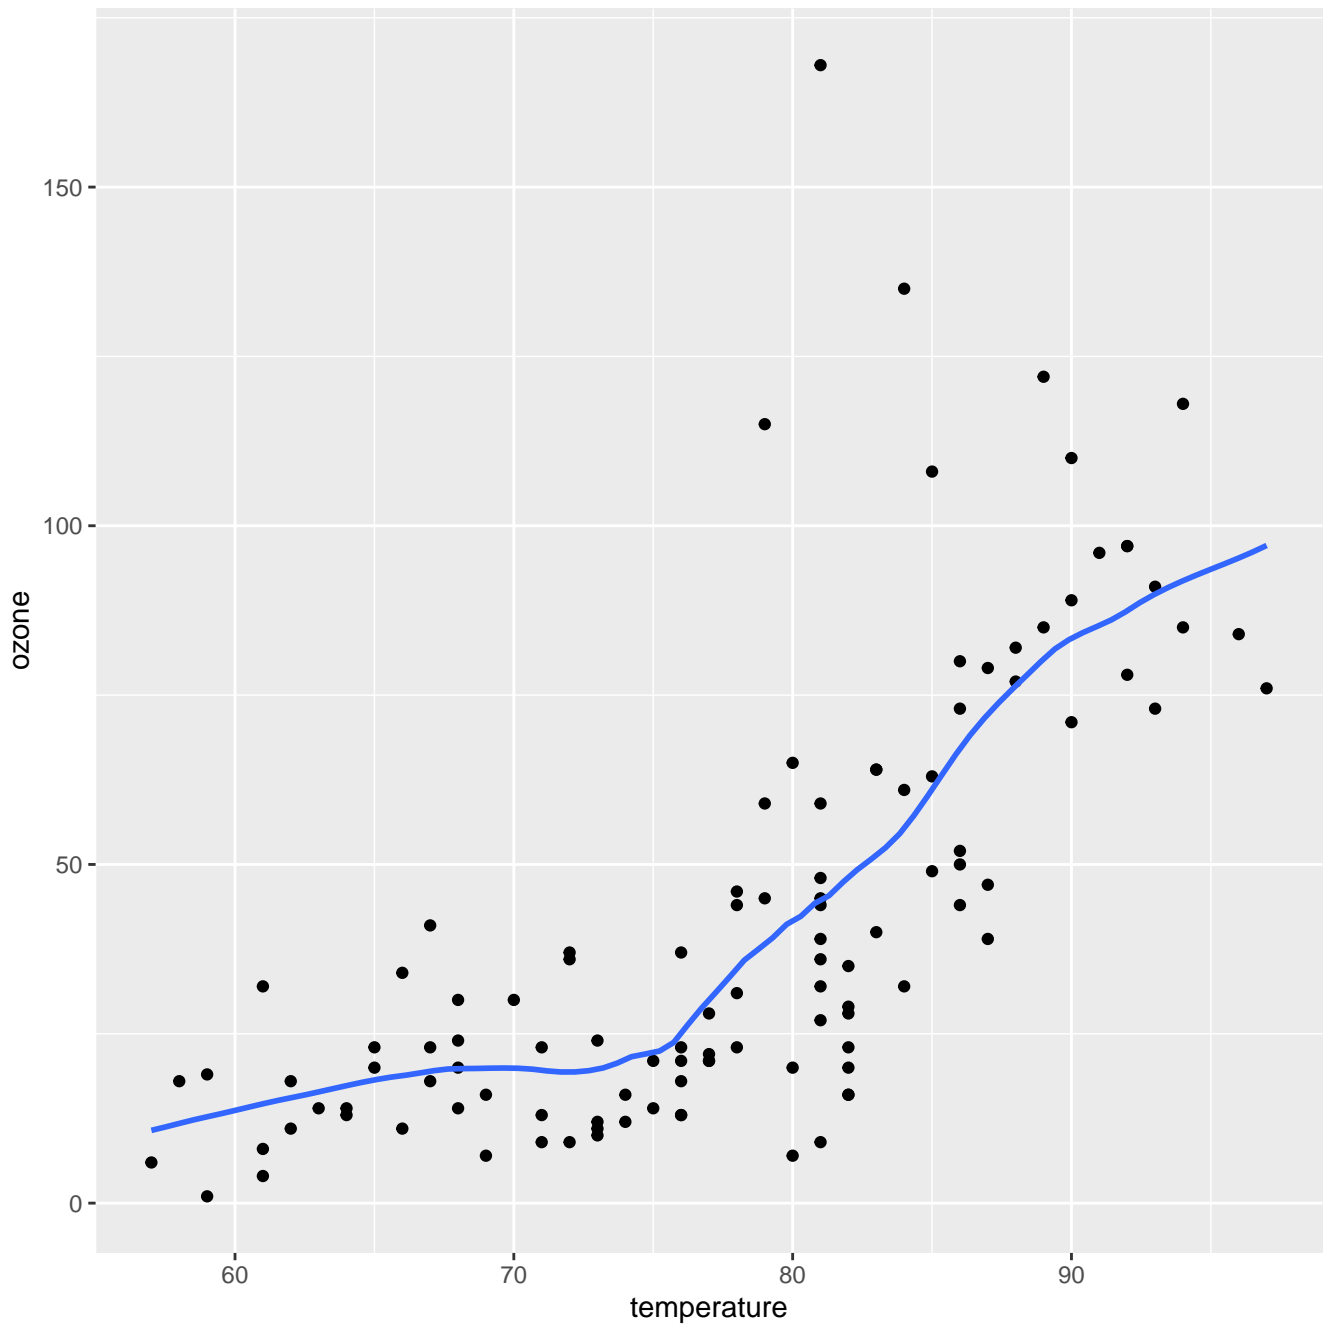


Visually inspecting the best fits compared to the 2nd and 3rd best fits, we do feel we may have over-fit the data, especially for the fits with degree 2. The model with degree 1 and span 0.4 appears to be the “best” fit.

Part c

Comparing results with built-in LOESS function

```
ozone %>% ggplot(aes(temperature, ozone)) +  
  geom_point() +  
  geom_smooth(method = "loess", degree = 1, span = 0.40, se = F, method.args = list(degree=1))
```

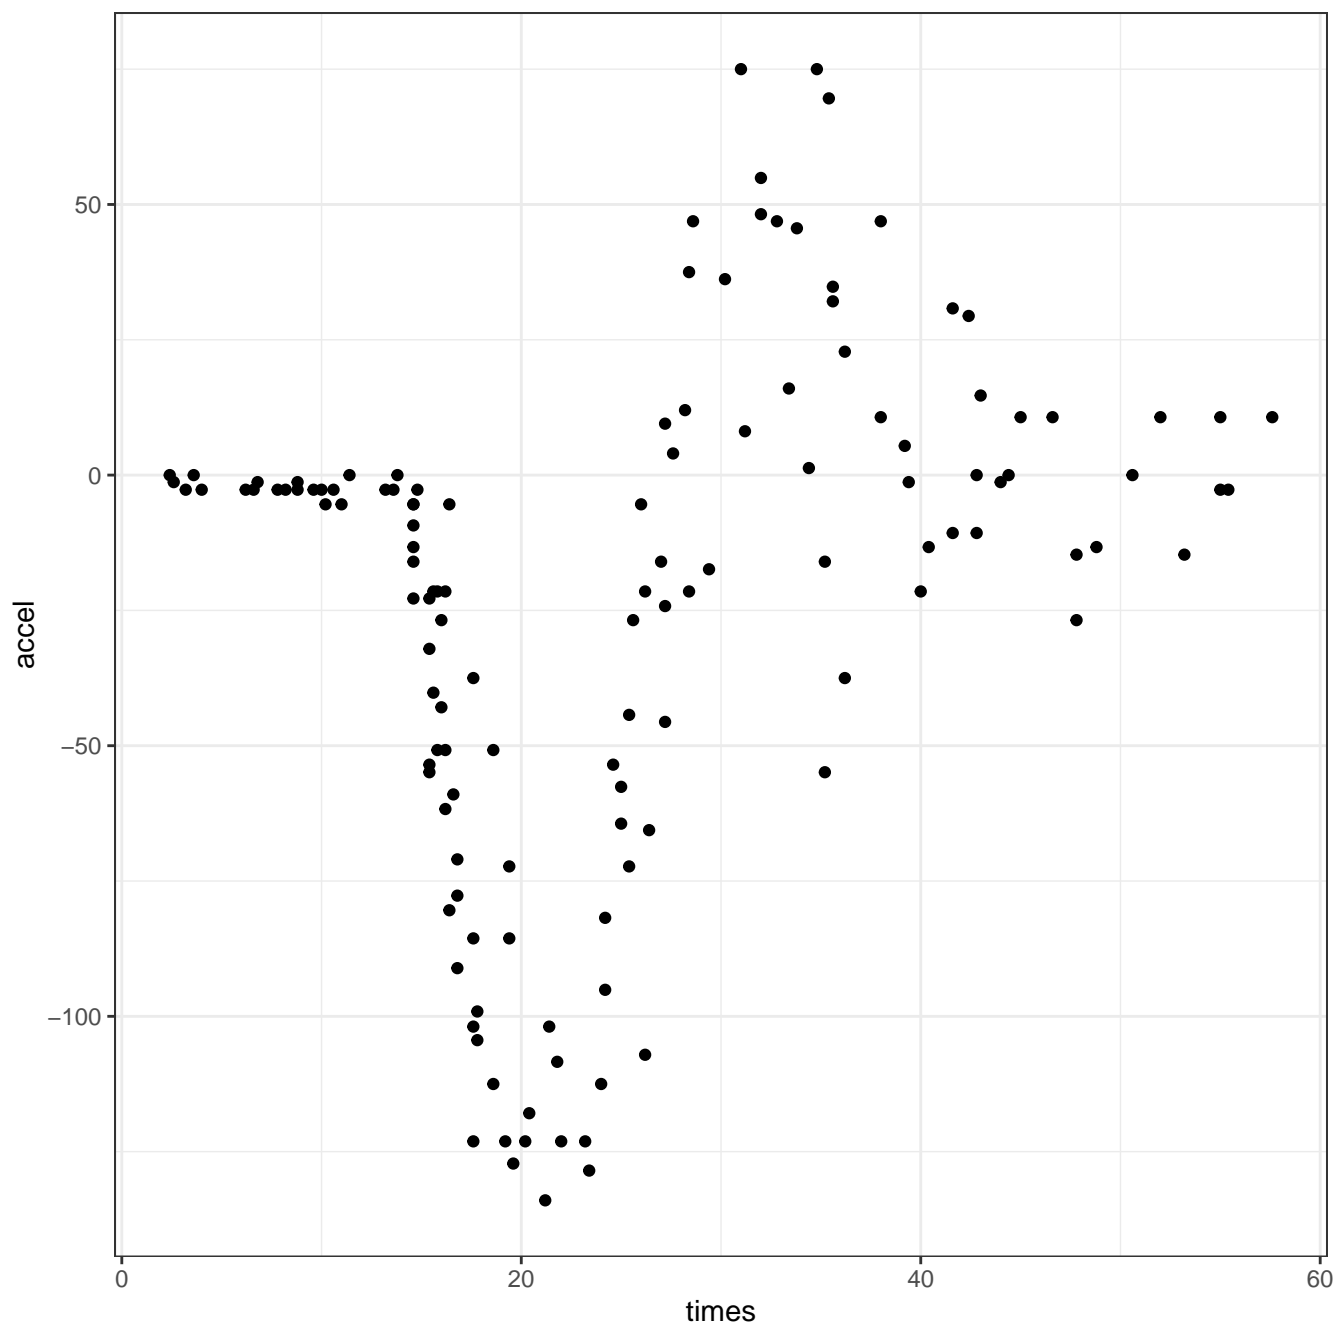


The built-in LOESS function results in a fit that does not over-fit the data, but we feel our model provides the better fit.

Problem 2

```
library(MASS)
data("mcycle")

ggplot(mcycle, aes(x = times, y = accel)) + theme_bw() + geom_point()
```



Part a

Determining LOESS regression fits on the data

```
# Creating an empty data frame
fit_table2 <- data.frame()

# Determining fits and putting info into data frame
for(j in 1:2) {
  for (i in seq(0.25, 0.75, by = 0.05)) {
    fit_table2 <- rbind(fit_table2, c(i, j, myloess(mcycle$times, mcycle$accel,
                                                    span = i, degree = j, show.plot = F)$RSE))
  }
}

# Changing column names
colnames(fit_table2) <- c("Span", "Degree", "RSE")
```

```
# Displaying data frame
fit_table2
```

	Span	Degree	RSE
1	0.25	1	24.54703
2	0.30	1	26.63732
3	0.35	1	28.88638
4	0.40	1	30.88313
5	0.45	1	32.57616
6	0.50	1	34.02896
7	0.55	1	35.27381
8	0.60	1	36.36153
9	0.65	1	37.32657
10	0.70	1	38.18843
11	0.75	1	38.96106
12	0.25	2	21.66747
13	0.30	2	21.99463
14	0.35	2	22.42284
15	0.40	2	23.12502
16	0.45	2	24.29479
17	0.50	2	25.87832
18	0.55	2	27.50811
19	0.60	2	29.01310
20	0.65	2	30.39499
21	0.70	2	31.68249
22	0.75	2	32.87589

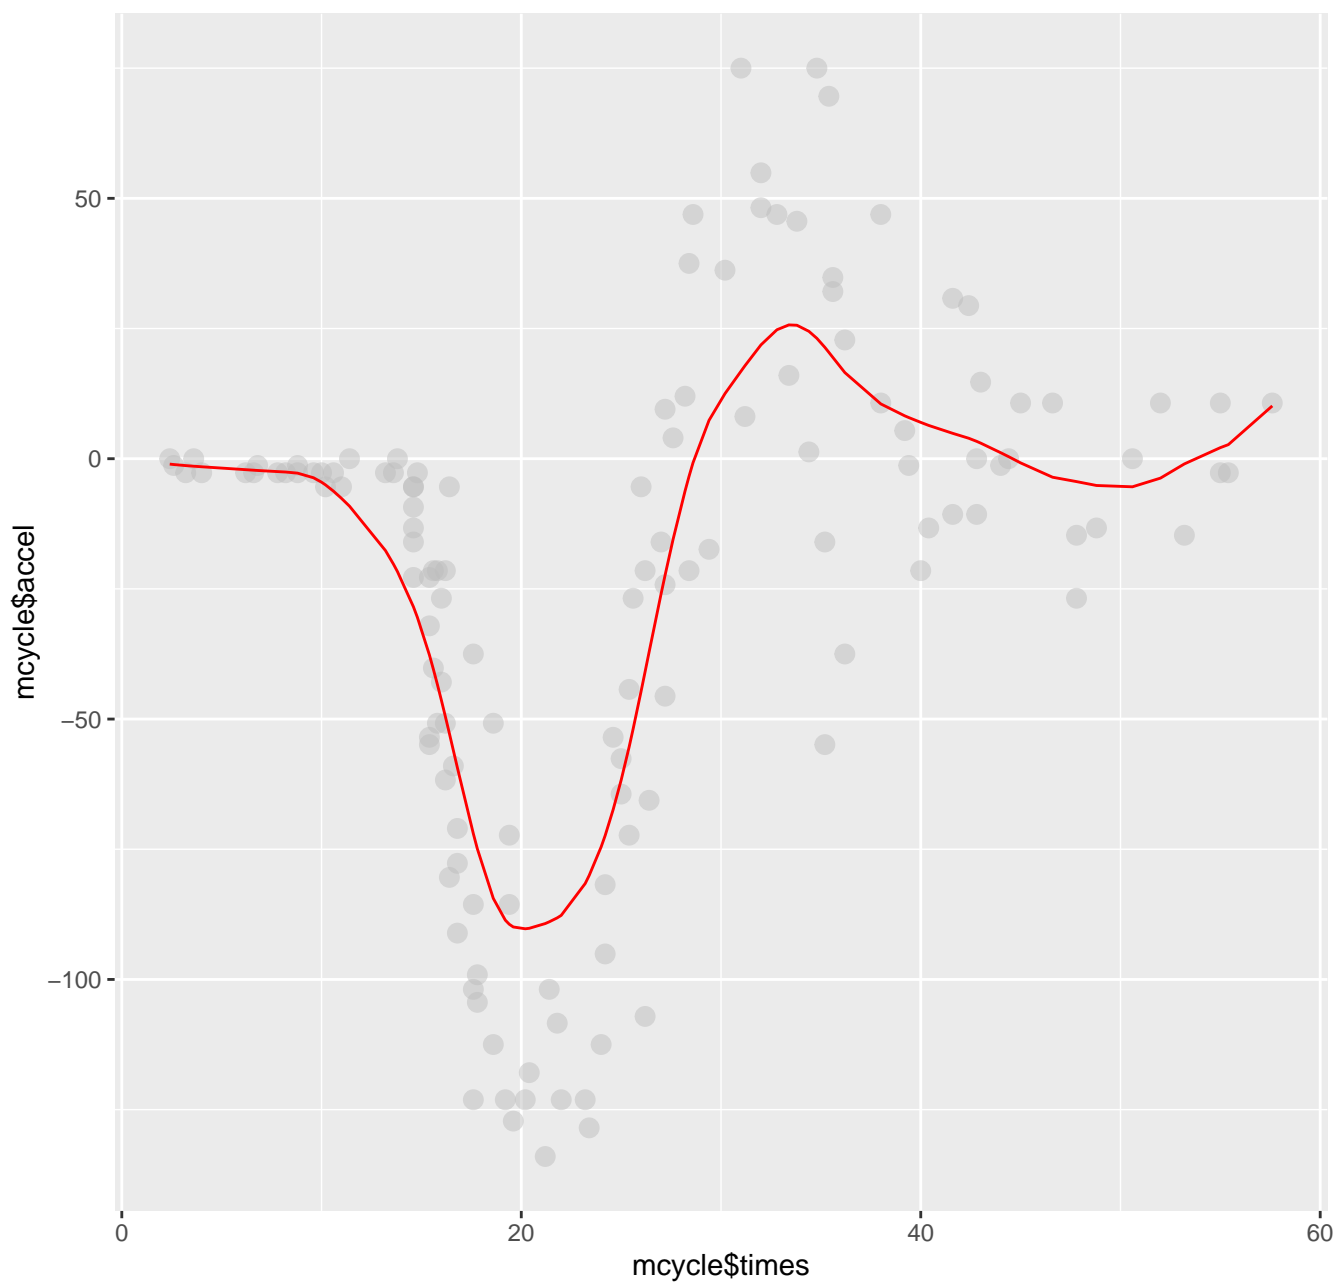
The three “best” fits with degree 1 appear to be span 0.25, 0.30, and 0.35. The three “best” fits with degree 2 appear to be span 0.25, 0.30, and 0.35.

Plotting the best fits found above

```
myloess(mcycle$times, mcycle$accel,
```

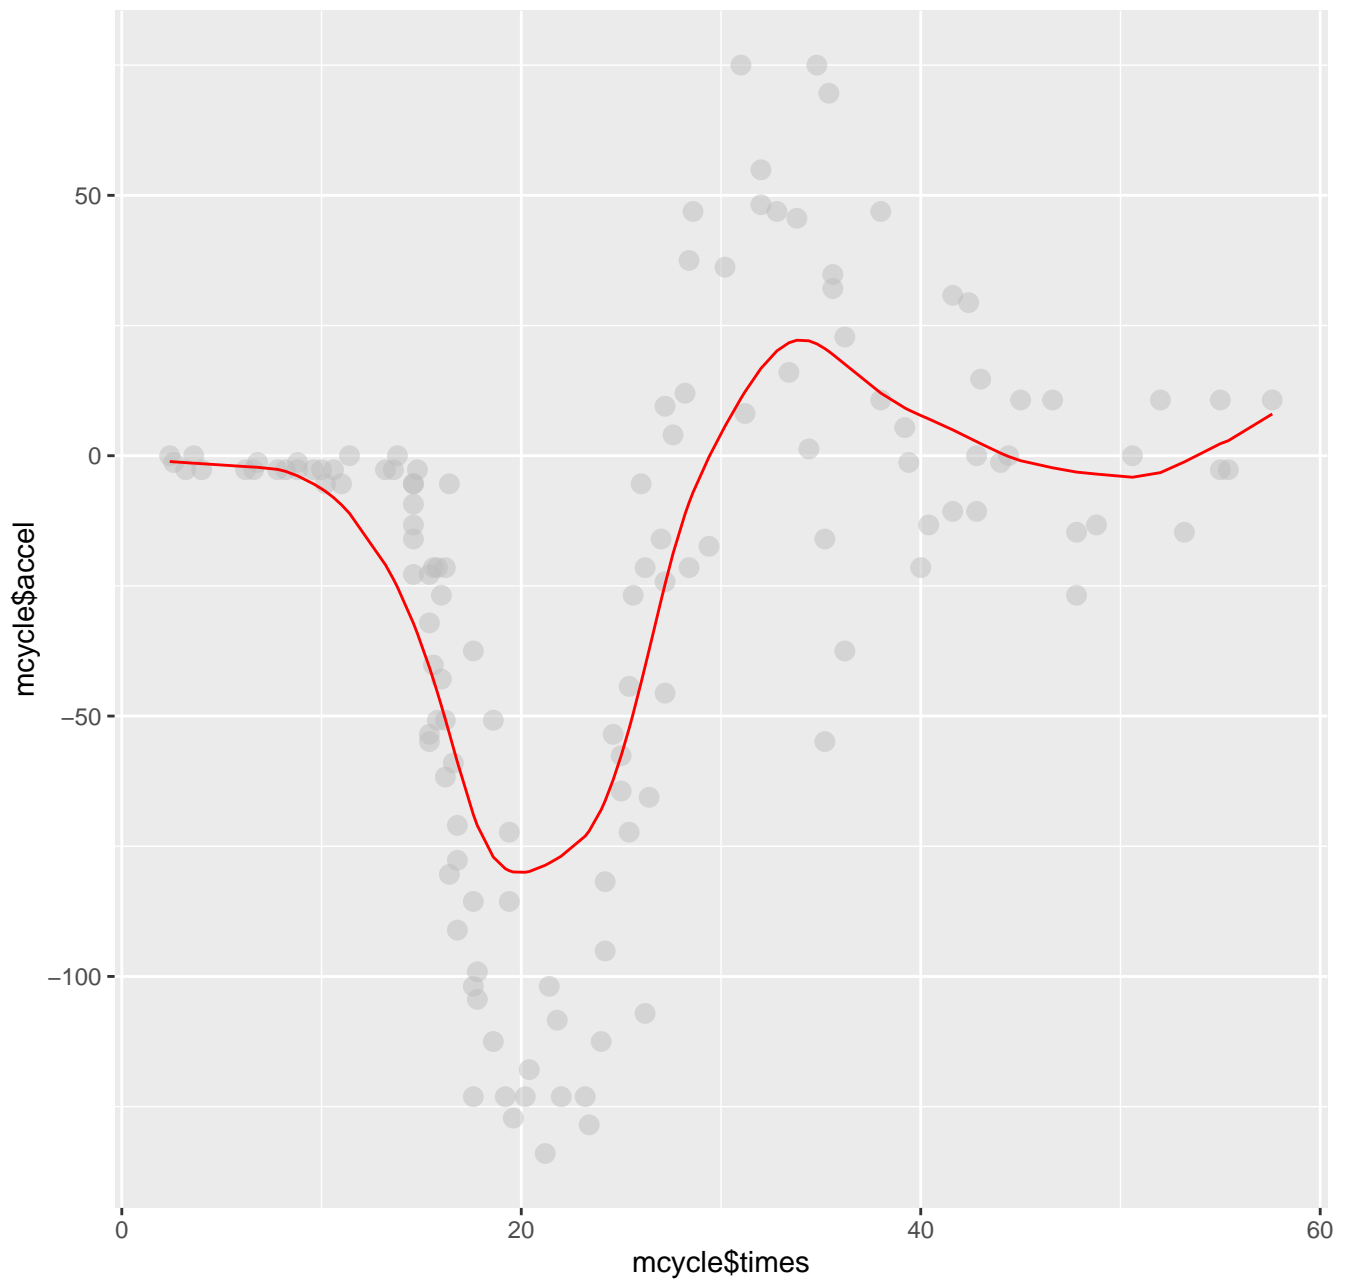
```
span = 0.25, degree = 1, show.plot = F)$loessplot
```


LOESS with degree = 1 and span = 0.25



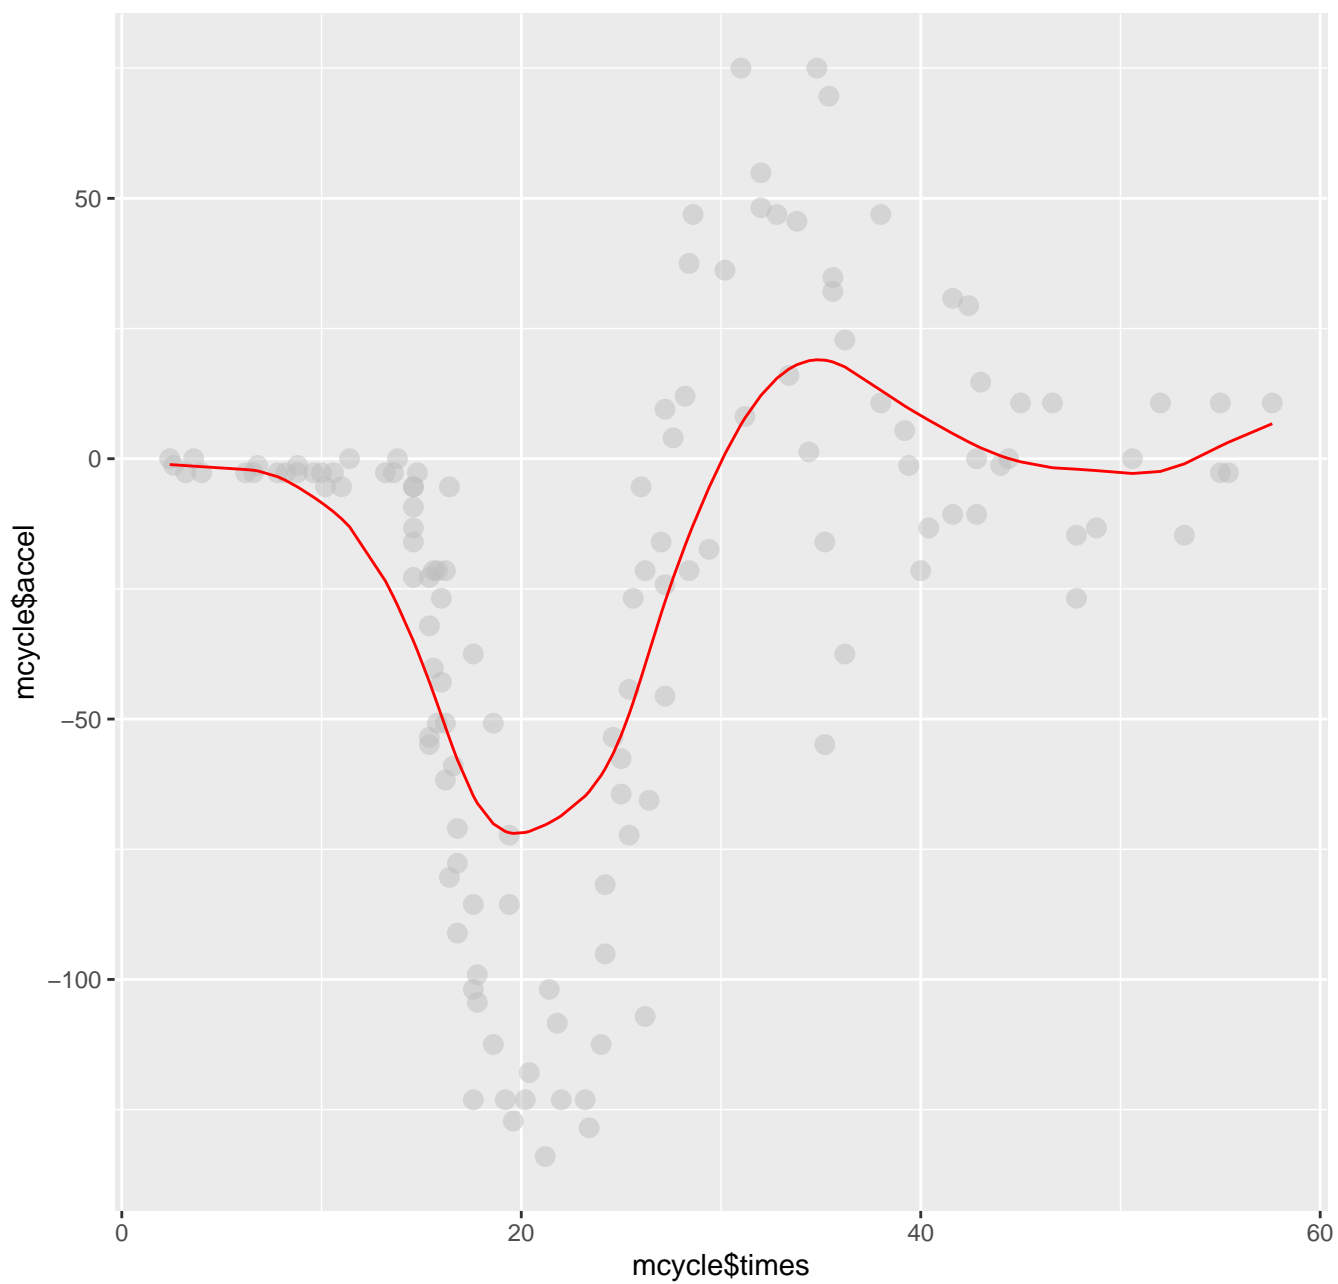
```
myloess(mcycle$times, mcycle$accel,  
span = 0.30, degree = 1, show.plot = F)$loessplot
```

LOESS with degree = 1 and span = 0.3



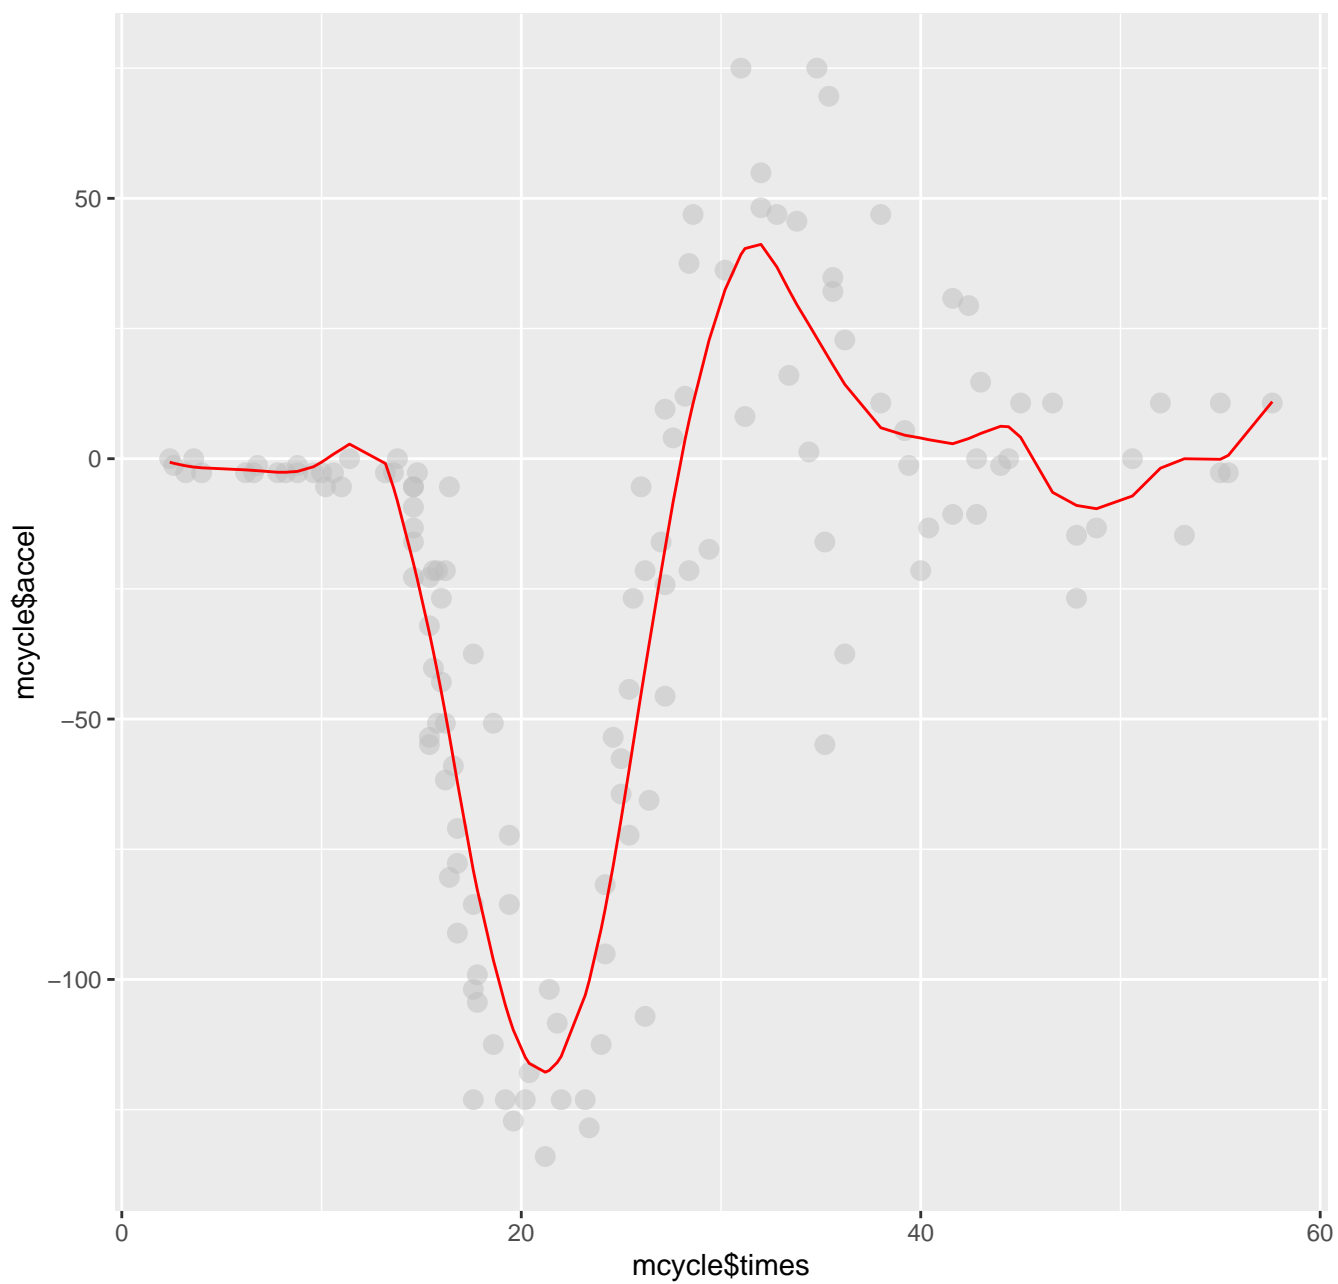
```
myloess(mcycle$times, mcycle$accel,  
span = 0.35, degree = 1, show.plot = F)$loessplot
```

LOESS with degree = 1 and span = 0.35



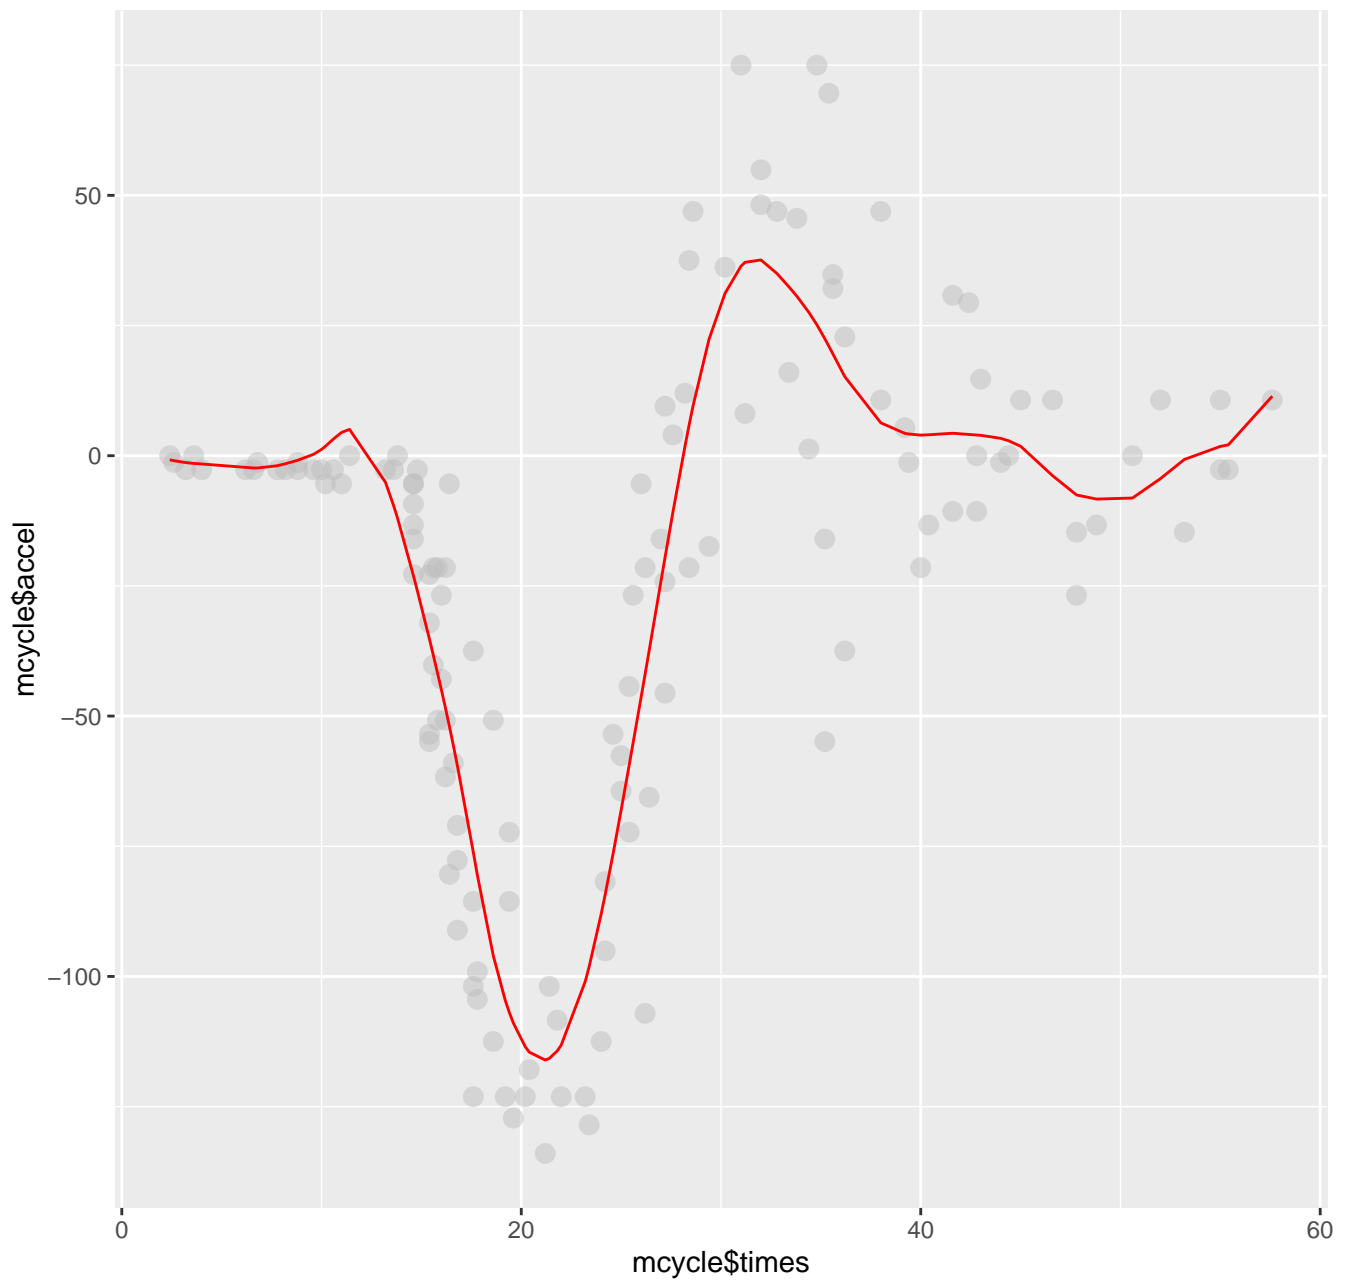
```
myloess(mcycle$times, mcycle$accel,  
span = 0.25, degree = 2, show.plot = F)$loessplot
```

LOESS with degree = 2 and span = 0.25



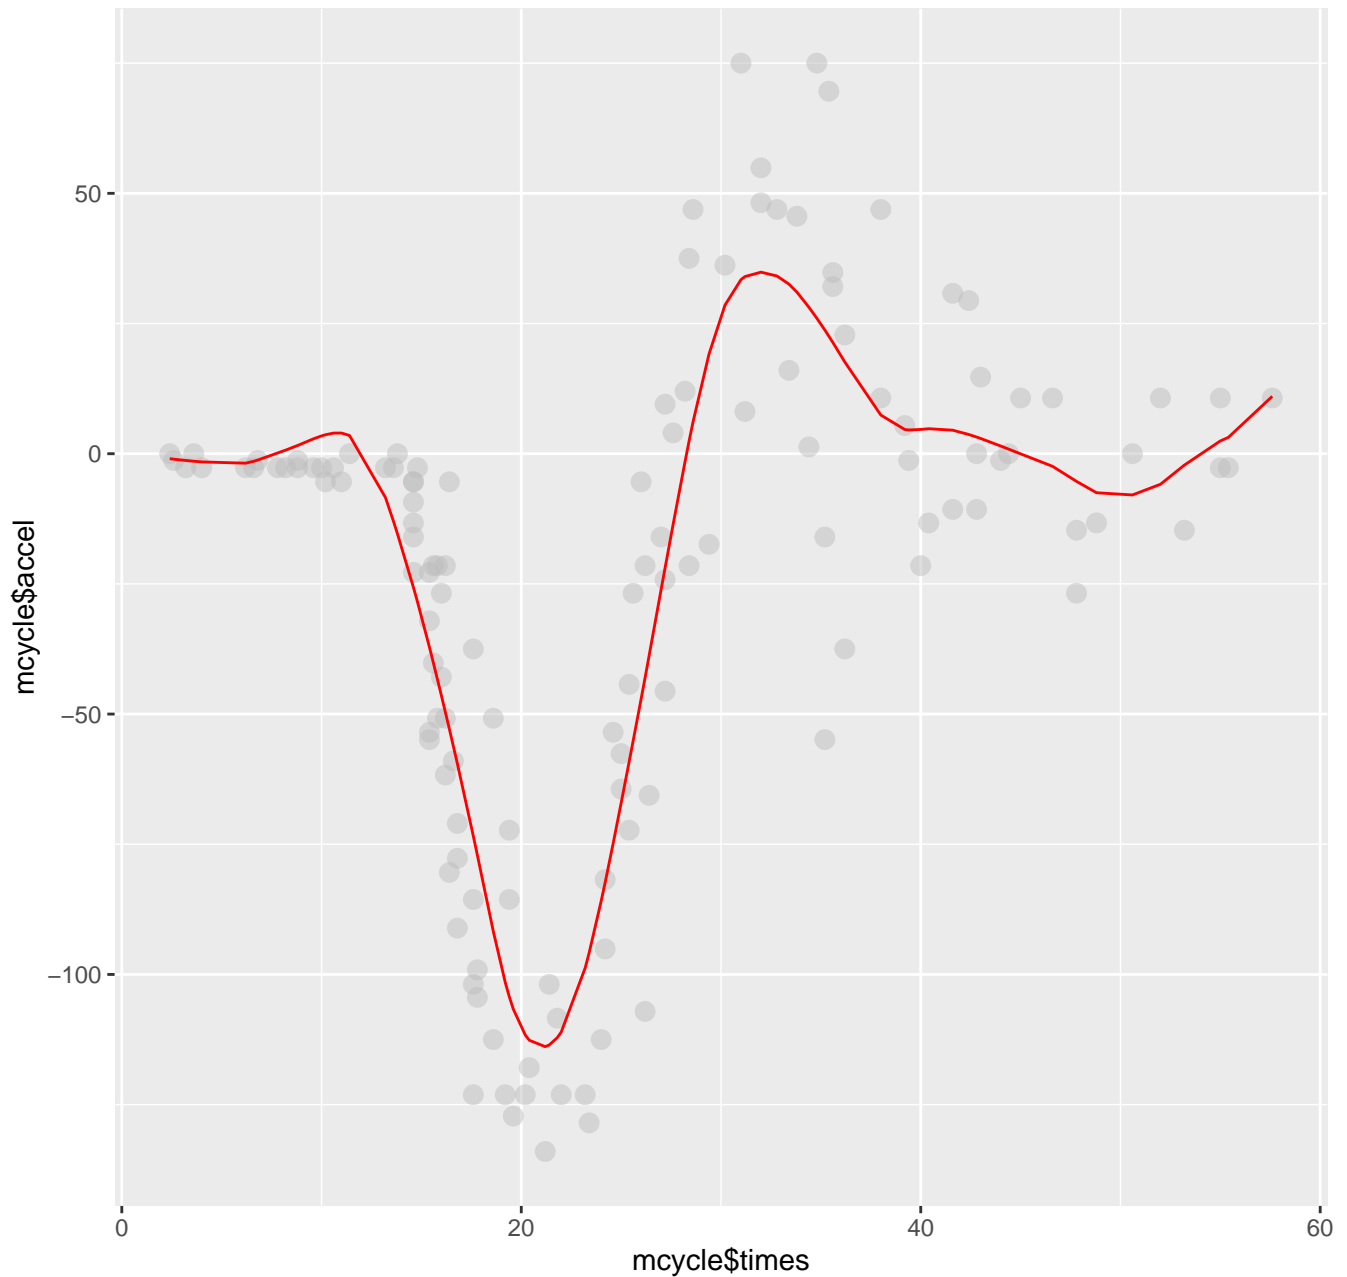
```
myloess(mcycle$times, mcycle$accel,  
span = 0.30, degree = 2, show.plot = F)$loessplot
```

LOESS with degree = 2 and span = 0.3



```
myloess(mcycle$times, mcycle$accel,  
span = 0.35, degree = 2, show.plot = F)$loessplot
```

LOESS with degree = 2 and span = 0.35

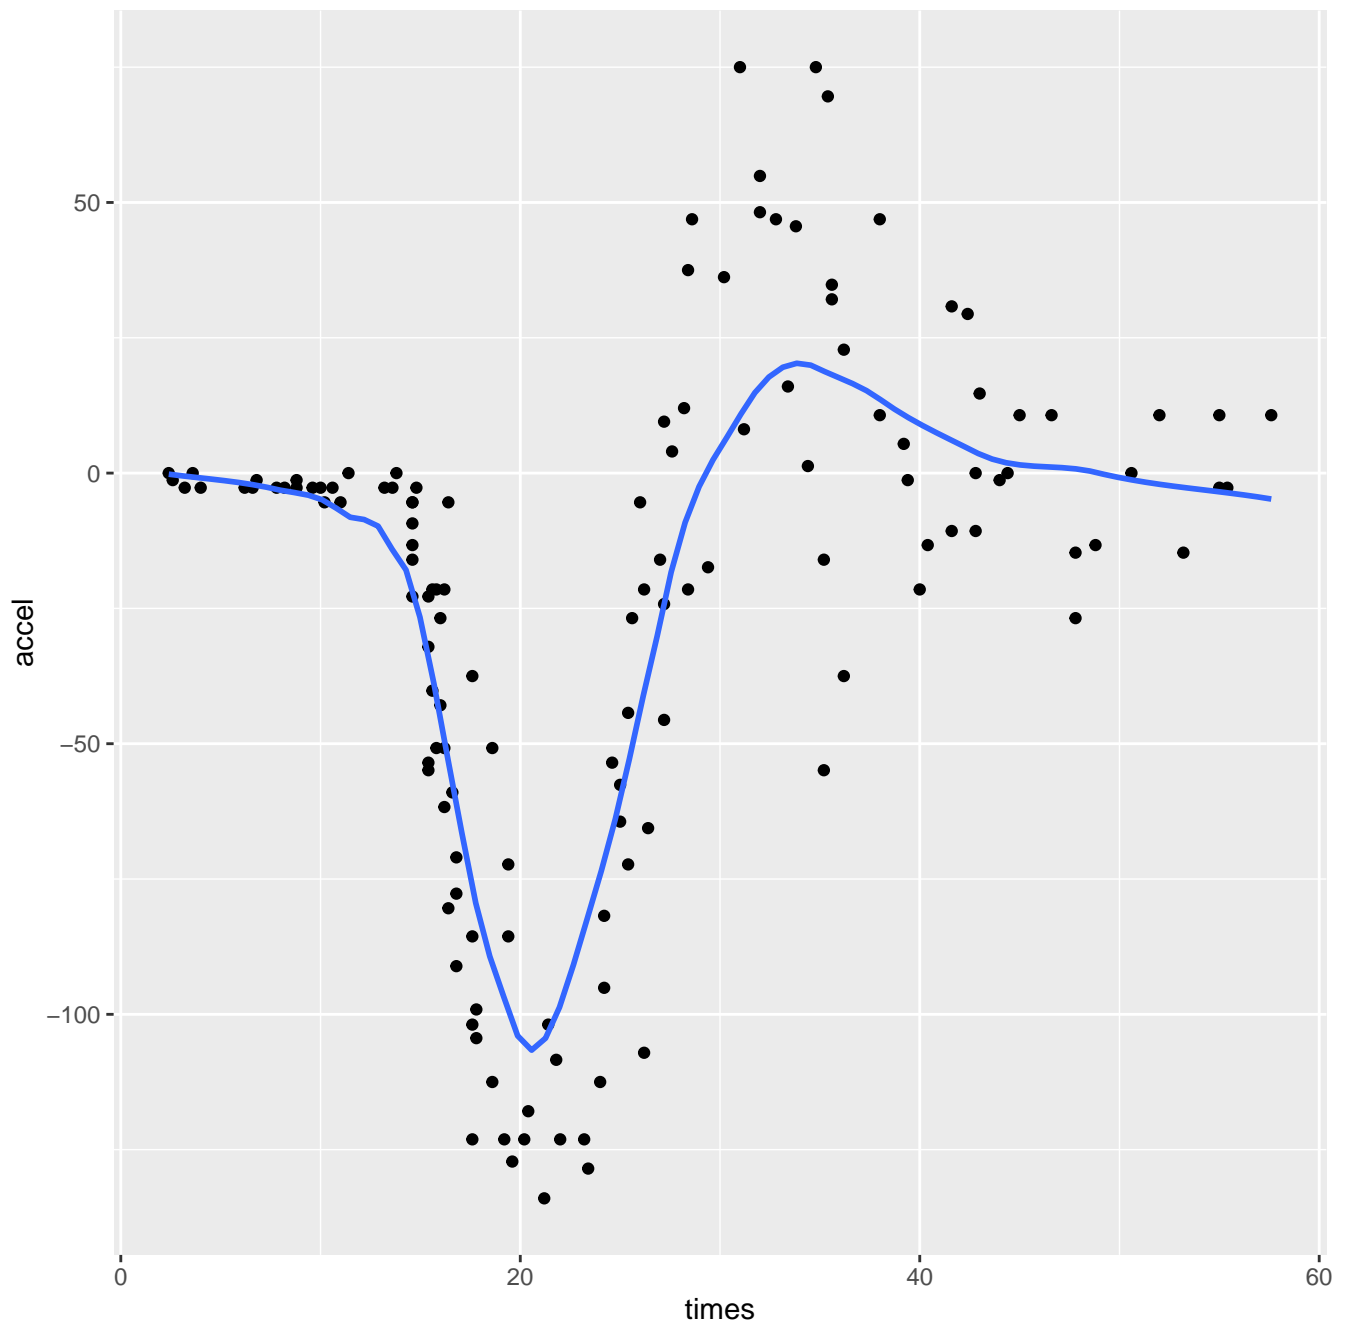


Visually inspecting the models, we believe that the model with degree 2 and span 0.35 provided the “best” fit.

Part b

Comparing results with built-in LOESS function

```
mcycle %>% ggplot(aes(times, accel)) +  
  geom_point() +  
  geom_smooth(method = "loess", degree = 2, span = 0.35, se = F, method.args = list(degree=1))
```



The built-in LOESS function results in a fit that does not over-fit the data, but we feel our model provides the better fit.

Problem 3

```
# Some pre-processing
library(ISLR)
# Remove the name of the car model and change the origin to categorical with actual name
Auto_new <- Auto[, -9]
# Lookup table
newOrigin <- c("USA", "European", "Japanese")
Auto_new$origin <- factor(newOrigin[Auto_new$origin], newOrigin)

# Look at the first 6 observations to see the final version
head(Auto_new)
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin
1	18	8	307	130	3504	12.0	70	USA

2	15	8	350	165	3693	11.5	70	USA
3	18	8	318	150	3436	11.0	70	USA
4	16	8	304	150	3433	12.0	70	USA
5	17	8	302	140	3449	10.5	70	USA
6	15	8	429	198	4341	10.0	70	USA